

Modern theory of detection and estimation

Lab 3: Linear filtering

Miguel Lázaro Gredilla

Introduction

In this lab session you will perform Bayesian and non-Bayesian filtering, first on toy data to grasp the basic concepts, and then on a real echo-cancellation scenario.

1. Working with toy data

We will consider first a toy data scenario. You will be generating observed data yourself and then trying to estimate the filter you used.

Use this code at the beginning of your script:

```
1 var_s = 1;  
2 var_n = 0.2;  
3 N = 100;  
4  
5 s = [0.54 1.83 -2.26 0.86 0.32]';  
6 randn('seed',0);rand('seed',0);
```

This sets the variance of the filter weights σ_s^2 , the variance of the noise σ_ϵ^2 , the length of the input and output signals, the filter **s** itself, and resets the random number generator, so that hopefully, despite using random numbers in this section, everyone gets the same results.

- 1.a) Start by generating an input signal $u[n]$, $0 \leq n < N$ by creating a random vector of length N containing independent realizations from the pdf $\mathcal{N}(0, 1)$. Filter it with **s** and add white Gaussian noise of variance **var_n** to produce signal **x**, which will also be of length N .

In order to create matrix \mathbf{U} you might want to check MatLAB's function `toeplitz`.

You can check that the output of the filter (before adding the Gaussian noise) is the same as the one obtained with MatLAB's function `filter`. Though you could use `filter` at this point and avoid building matrix \mathbf{U} altogether, you will be needing it later, so it's recommended that you only use the `filter` command as a correctness check.

- 1.b) Compute the predictive probability density of the filter given the observed output $p(\mathbf{s}|\mathbf{x})$. You should be getting something like this:

```

1 mean_s =
2
3     0.5204
4     1.7575
5    -2.3716
6     0.8857
7     0.3142
8
9
10 cov_s =
11
12     0.0022     0.0002     0.0001     0.0002    -0.0001
13     0.0002     0.0023     0.0002     0.0002     0.0002
14     0.0001     0.0002     0.0023     0.0002     0.0002
15     0.0002     0.0002     0.0002     0.0023     0.0002
16    -0.0001     0.0002     0.0002     0.0002     0.0023

```

Provide an interpretation for this result.

- 1.c) Let's turn now to prediction. If the next input sample, $u[100]$, was equal to zero, what would the predictive probability density of the corresponding output $x[100]$ (after traversing the filter and i.i.d. Gaussian noise is added) be? Compare your result with the correct one:

```

1 mu_star =
2     2.3322
3
4 v_star =
5     0.2036

```

2. Echo cancellation scenario

In this section we consider a scenario of voice chat over a noisy channel with echo.

You will find two appropriately named variables in the attached data file `data.mat`. These are the voice registers, sampled at 22.05 KHz and containing 100 Ksamples each. The provided vectors include data for $0 \leq n \leq N - 1$, with $N = 100000$.

Variable `voiceout` contains the voice of the local speaker as recorded on the local machine, whereas variable `voicein` contains the voice of the remote speaker as received on the local machine plus a distorted version of the local speaker's voice after traversing the channel to the remote machine and back (out-going and incoming channels are not necessarily identical).

More formally:

$$\text{voicein}[n] = \text{voiceremote}[n] + \sum_{k=-\infty}^{\infty} \text{voiceout}[n - k] s_n[k] \quad (1)$$

where $s_n[k]$ is a causal, time-varying (hence the subscript n) FIR filter that models the full round trip of the local voice. Signal `voiceremote` includes any distortion to the original remote speaker's voice that is introduced by the incoming channel, as well as any additional noise — we will not attempt to revert those, but only remove distortions related to `voiceout`.

MatLAB provides functions that will allow you to play these files (as well as any reconstruction of `voiceremote` that you may have). You can use this to informally test the quality of your reconstruction.

The round-trip channel model $s_n[k]$ is known to have very few non-zero values. More specifically, $s_n[k]$ is non-zero only for the following values of k :

```
1 nonzerolags = [3122 5953 9999 14999 18999 29295 39385];
```

From a programmatic point of view, be reminded that, in mathematical notation, $s_n[k]$ is a zero-based vector, whereas MatLAB uses one-based vectors.

- 2.a) For this part, disregard the time variation of the round-trip channel, and model it as a constant round-trip channel $s[k]$.

Provide a probabilistic estimation of $s[k]$. You know that only a few values are non-zero, so focus on those elements only. You can use the following ground-truth values in this case:

```
1 var_s = 0.2;
2 var_n = 4e-5;
```

Pitfall: If you try to use the standard equations from the notes in this case, you will run into trouble twice. First, when you try to compute \mathbf{U} , you will, almost surely, run out of memory. Second, even if you used a computer from the future, applying the standard equations to compute the probabilistic prediction will result in placing non-zero probability on non-zero values for all coefficients of $s[k]$. And we know that most coefficients are exactly zero, so that belief is incorrect.

Hint: Follow the derivations in the notes, but use the fact that only a few values on $s[k]$ are non-zero to obtain a very similar, yet much faster equation to compute a solution that provides a joint posterior for the non-zero coefficients (and by definition predicts exactly zero for all the remaining coefficients).

You should be getting something like this:

```
1 mean_s =
2
3     1.3189
4     0.5548
5     0.5585
6     0.0007
7     0.2646
8     0.2187
9     0.2594
10
11
12 cov_s =
13
14     1.0e-06 *
15
16     0.2043  -0.0087   0.0001  -0.0000  -0.0031   ...
17           0.0002  -0.0000
18     -0.0087   0.2042  -0.0001   0.0002   0.0006   ...
19           0.0018  -0.0000
```

18	0.0001	-0.0001	0.2040	0.0003	0.0005	...
	0.0053	0.0000				
19	-0.0000	0.0002	0.0003	0.2039	0.0014	...
	-0.0003	0.0024				
20	-0.0031	0.0006	0.0005	0.0014	0.2039	...
	0.0001	0.0002				
21	0.0002	0.0018	0.0053	-0.0003	0.0001	...
	0.2040	-0.0002				
22	-0.0000	-0.0000	0.0000	0.0024	0.0002	...
	-0.0002	0.2039				

- 2.b) Provide an MMSE estimation of **voiceremote**. Listen to the provided **voicein** and **voiceout** and then to your obtained estimation of **voiceremote**.

When listening to your estimation, you should be able to hear some echos of **voiceout** and someone else talking faintly that could not be heard when listening to **voicein** or **voiceout**.

- 2.c) Now make use the LMS algorithm. As mentioned in class, this algorithm enjoys adaptivity properties, so it should be able to *track* the channel, which is actually time-varying. Remember that $s_n[k]$ is non-zero only for the values of k mentioned before.

Also, note that the LMS algorithm provides an online estimation of **voiceremote**. I.e., after you process sample n from **voiceout** and **voicein**, you already get an estimation of sample n of **voiceremote**. This means that estimations can be provided on-the-fly, instead of in batch form as we did previously.

Hint: Using the vanilla LMS algorithm will result in unnecessary overhead. You can simplify it by taking advantage of having just a handful of non-zero coefficients. This is not required, but will turn out to be convenient.

If you listen to **voiceremote_LMS**, you should be able to hear a distorted version of the remote speaker without clear echos of **voiceout**.

- 2.d) You can check that this result is better than the previous one just by listening to it and providing your subjective impression. Can you think of a way to provide an objective assessment of the quality of this estimation as compared to your previous estimation?

3. Extension exercise

3.a) The evolution of every non-zero coefficient at lag k follows this equation

$$s_n[k] = c^{(k)} + \sum_{j=1}^{10} a_j^{(k)} \cos(2\pi n j / f_s + \phi_j^{(k)}).$$

with $f_s = 22050$ Hz. I.e., each coefficient is the composition of a constant term plus up to 10 sinusoids. We say “up to”, because we can remove any or even all sinusoids for a given coefficient by setting the corresponding $a_j^{(k)}$ ’s to zero.

Estimate the values of $\{c^{(k)}, a_j^{(k)}, \phi_j^{(k)}\}$ for every k corresponding to a non-zero coefficient and $1 \leq j \leq 10$.

Hint: Cast the problem as a linear regression problem.

3.b) Use the previous estimation of $s_n[k]$ to estimate **voiceremote**. If you listen to it you should hear an almost perfect reconstruction: No echos, no sound distortion.

4. Reference implementation

As a reference, a complete implementation of the solution to this questions took less than 70 lines of MatLAB code (including comments and proper spacing) and takes less than 10 seconds to run on a mid-2012 macbook air. Part 3.a) is the most computationally intense and can take up to 1GB of memory.